

南京大学本科生实验报告

课程名称：计算机网络

任课教师：李文中

助教：

学院	计算机科学与技术系	专业（方向）	计算机科学与技术
学号	211220049	姓名	石璐
Email	211220049@smail.nju.edu.cn	开始/完成日期	2023 年 5 月 4 日

1. 实验名称：Forwarding Packets

2. 实验目的

接收和转发到达链路并以其他主机为目的地的数据包。转发过程的一部分是在转发表中进行地址查找（"最长前缀匹配"查找）。

为没有已知以太网 MAC 地址的 IP 地址发出 ARP 请求。路由器经常需要向其他主机发送数据包，并需要以太网 MAC 地址来实现。

3. 实验内容

Task 1: IP Forwarding Table Lookup

✓ Coding:

1) Build Forwarding Table

- 定义转发表项

```
12 class ForwardingTableItem:
13     def __init__(self,ipaddr,mask,nexthop,interface):
14         self.ipaddr = IPv4Address(ipaddr)
15         self.mask = IPv4Address(mask)
16         self.nexthop = IPv4Address(nexthop)
17         self.interface = interface
```

- 建立转发表

(注: 按前缀长度降序排序, 便于"最长前缀匹配"查找)

```
23 class ForwardingTable:
24     def __init__(self,interfaces):
25         # from router interfaces
26         self.table = [ForwardingTableItem(intf.ipaddr,intf.netmask,'0.0.0.0',intf.name) \
27                       for intf in interfaces]
28         # from file
29         with open('forwarding_table.txt') as f:
30             for line in f.readlines():
31                 line = line.strip('\n')
32                 ipaddr,mask,nexthop,interface = line.split(' ')
33                 self.table.append(ForwardingTableItem(ipaddr,mask,nexthop,interface))
34         # sort forwarding table
35         self.table.sort(key=prefix_len,reverse=True)
```

2) Match Destination IP Addresses against Forwarding Table

由于转发表已经按前缀长度降序排序, 因此顺序遍历的第一个匹配结果即为"最长前缀匹配"查找的结果。

```
23 class ForwardingTable:
24 > def __init__(self,interfaces): ...
41
42     def search(self,ipaddr):
43         for item in self.table:
44             prefixnet = IPv4Network(f'{item.ipaddr}/{item.mask}',strict=False)
45             if ipaddr in prefixnet:
46                 return item
47         return None
```

Task 2: Forwarding the Packet and ARP

✓ Coding:

- 等待队列 waitingtable

采用字典结构, key 为未知 mac 地址的 IP 地址, value 为 waitinglist 类, 类中存储了 4 项信息:

- a) **time:** 针对该 ip 地址最近一次发送 ARP request 包的时间
- b) **trycnt:** 剩余可发送 request 的次数, 初始为 5 次, 减为 0 时将丢弃队列中等待的 packets
- c) **intf:** 对应的路由器端口
- d) **packets:** 等待发送到该 ip 的 packet 队列, 保序

```
56 class Router(object):
57     def __init__(self, net: switchyard.llnetbase.LLNetBase):
58         ...
59         self.waitingtable = {}
```

```
49 class WaitingList:
50     def __init__(self, intf):
51         self.time = time.time()
52         self.trycnt = 5
53         self.intf = intf
54         self.packets = []
```

- 处理等待队列

每次通过主循环时处理队列中的项目, 看看是否需要发送 ARP 请求重传或丢弃等待的 packets。

```
165 def start(self):
166     '''A running daemon of the router.
167     Receive packets until the end of time.
168     '''
169     while True:
170         try:
171             self.handle_waitinglist()
172             recv = self.net.recv_packet(timeout=1.0)
```

```
149 def handle_waitinglist(self):
150     for ip, wlst in self.waitingtable.items():
151         if len(wlst.packets) > 0 and time.time() - wlst.time >= 1:
152             # Retransmission
153             if wlst.trycnt > 0:
154                 arp_request = create_ip_arp_request(wlst.intf.ethaddr, wlst.intf.ipaddr, ip)
155                 log_info(f"Sending packet {arp_request} to {wlst.intf.name}")
156                 self.net.send_packet(wlst.intf, arp_request)
157                 self.waitingtable[ip].time = time.time()
158                 self.waitingtable[ip].trycnt -= 1
159             # Drop packets
160             else:
161                 self.waitingtable[ip].trycnt = 5
162                 self.waitingtable[ip].packets.clear()
```

- 检查 Ethernet header

```
76         # Drop the packet with Ethernet destination neither a broadcast address
77         # nor the MAC of the incoming port.
78         if ethheader and ethheader.dst != "ff:ff:ff:ff:ff:ff" \
79             and ethheader.dst != self.net.interface_by_name(ifaceName).ethaddr:
80             return
81         # Drop the packet with VLAN
82         if ethheader.ethertype == EtherType.x8021Q:
83             return
```

- 处理 IPv4 packet

- 将 ttl 字段递减 1
- 在转发表中查找 ip 目标地址对应的表项
- 如果 nexthop 为 0.0.0.0，则下一跳的 ip 地址设置为表项 ipaddr 域
- 为要转发的 IP 数据包创建一个新的 Ethernet header。此时需要下一跳的 mac 地址，首先查询 arp 表，如果查到直接转发，如果未查到则将该数据包放入等待队列并发送 arp 请求报文。

```
113 if ipheader:
114     # Drop the packet for the router itself.
115     # TODO: handle in lab 5
116     if ipheader.dst not in self.ipaddrs:
117         packet[IPv4].ttl -= 1
118         # Process the packet not expired.
119         # TODO: handle in lab 5
120         if packet[IPv4].ttl > 0:
121             item = self.forwardingtable.search(packet[IPv4].dst)
122             # Drop the packet if there is no match in the table.
123             # TODO: handle in lab 5
124             if item:
125                 intf = self.net.interface_by_name(item.interface)
126                 # The IP address 0.0.0.0 means that if a packet's destination IP matches this n
127                 # just throw the packet out of this interface, expecting that the next hop is t
128                 dstip = item.nexthop
129                 if dstip == IPv4Address('0.0.0.0'):
130                     dstip = packet[IPv4].dst
```

```
131         # search ARP Table for dst mac
132         if dstip in self.arptable:
133             dstmac = self.arptable[dstip]
134             packet[0] = Ethernet(src=intf.ethaddr,dst=dstmac,ethertype=EtherType.IPv4)
135             log_info (f"Sending packet {packet} to {intf.name}")
136             self.net.send_packet([intf,packet])
137         # send ARP request for dst mac
138         else:
139             if dstip not in self.waitingtable:
140                 self.waitingtable[dstip] = WaitingList(intf)
141             if len(self.waitingtable[dstip].packets) == 0:
142                 arp_request = create_ip_arp_request(intf.ethaddr,intf.ipaddr,dstip)
143                 log_info (f"Sending packet {arp_request} to {intf.name}")
144                 self.net.send_packet(intf, arp_request)
145                 self.waitingtable[dstip].time = time.time()
146                 self.waitingtable[dstip].trycnt -= 1
147                 self.waitingtable[dstip].packets.append(packet)
```

- 处理 ARP packet

- i. 更新 ARP 表。
- ii. 处理 ARP request。发送 ARP 回复报文。
- iii. 处理 ARP reply。完成要转发的 IP 数据包的以太网头，并与等待队列中的 packets 一起发送。

```

84 # Only process the ARP packet when its destination IP is held by a port of the router.
85 if arpheader and arpheader.targetprotoaddr in self.ipaddrs:
86     # Cached ARP Table
87     if arpheader.operation == ArpOperation.Request or arpheader.senderhwaddr != "ff:ff:ff:ff:ff:ff":
88         self.arptable[arpheader.senderprotoaddr] = arpheader.senderhwaddr
89     # Handle ARP packet
90     if arpheader.targetprotoaddr in self.ipaddrs:
91         # Handle ARP request
92         if arpheader.operation == ArpOperation.Request:
93             intf = self.net.interface_by_name(ifaceName)
94             reply_mac = self.net.interface_by_ipaddr(arpheader.targetprotoaddr).ethaddr
95             arp_reply = create_ip_arp_reply(reply_mac, arpheader.senderhwaddr, \
96                                             arpheader.targetprotoaddr, arpheader.senderprotoaddr)
97             log_info(f"Sending packet {arp_reply} to {intf.name}")
98             self.net.send_packet(intf, arp_reply)
99         # Handle ARP reply
100     elif arpheader.operation == ArpOperation.Reply:
101         if arpheader.senderhwaddr != "ff:ff:ff:ff:ff:ff":
102             ipaddr, ethaddr = arpheader.senderprotoaddr, arpheader.senderhwaddr
103             if len(self.waitingtable[ipaddr].packets) > 0:
104                 intf = self.waitingtable[ipaddr].intf
105                 e = Ethernet(src=intf.ethaddr, dst=ethaddr, ethertype=EtherType.IPV4)
106                 for waiting_packet in self.waitingtable[ipaddr].packets:
107                     waiting_packet[0] = e
108                     log_info(f"Sending packet {waiting_packet} to {intf.name}")
109                     self.net.send_packet(intf, waiting_packet)
110                 self.waitingtable[ipaddr].packets.clear()
111                 self.waitingtable[ipaddr].trycnt = 5

```

 Testing:

```

1203Ping request from 31.0.6.1 should arrive on eth6
1204Ping request from 31.0.6.1 should arrive on eth6
1205Router should not do anything
1206Bonus: V2FybSB1cA==
1207Bonus: Q29vbCBkb3du
1208Bonus: V2hldCBkcyB5YSBob3BlIHQnIGZpbmQgaGVyZT8=

Failed:
    Bonus: Tm90aGluJyBmb3IgewEgdCcgZmluZCBoZXJlIQ==
        Expected event: recv packet Ethernet
        30:00:00:00:05:02->20:00:00:00:00:05 IP | IPv4
        15.0.0.92->192.168.129.2 UDP | UDP 0->0 RawPacketContents
        (35 bytes) b'\x00\x00\x00\x00\x00\x00\x00\x00\x00\x00...' on
        eth5

Pending (couldn't test because of prior failure):
1   Q29uZ3JhdHMh

*****
Your code didn't crash, but a test failed.
*****

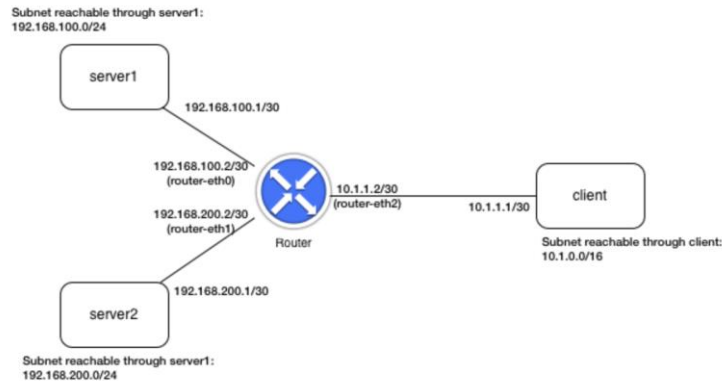
This is the Switchyard equivalent of the blue screen of death.
As far as I can tell, here's what happened:

Expected event:
    Bonus: Tm90aGluJyBmb3IgewEgdCcgZmluZCBoZXJlIQ==

Failure observed:
    send packet was called instead of recv packet Ethernet
    30:00:00:00:05:02->20:00:00:00:00:05 IP | IPv4
    15.0.0.92->192.168.129.2 UDP | UDP 0->0 RawPacketContents
    (35 bytes) b'\x00\x00\x00\x00\x00\x00\x00\x00\x00\x00...' on
    eth5

```


✅ Deploying: 执行 server1# ping -c2 10.1.1.1



- server1(192.168.100.1)欲向 client 发送 ICMP 包, 首先要向 router 的 eth0 端口发送, 因此 server1 询问 eth0 端口(192.168.100.2)的 mac 地址, 在 eth0 收到一条 192.168.100.1 对 192.168.100.2 的 ARP 请求报文。
- router 回复上述 ARP 请求, 在 eth0 发送一条 192.168.100.2 对 192.168.100.1 的 ARP 回复报文。
- server1 得到 eth0 的 mac 地址后开始发送 ICMP 报文, 因此在 eth0 收到一条 192.168.100.1 对 10.1.1.1 的 ICMP 报文。
- router 查询转发表后将该报文导向 eth2 端口准备发向 client, 但是要先询问 10.1.1.1 的 mac 地址, 因此在 eth2 发送一条 10.1.1.2 对 10.1.1.1 的 ARP 请求报文, 同时将 ICMP 报文缓存起来。
- client 回复上述 ARP 请求, 在 eth2 收到一条 10.1.1.1 对 10.1.1.2 的 ARP 回复报文。
- router 得到 client 的 mac 地址后发送之前缓存的 ICMP 报文, 因此在 eth2 发送一条 192.168.100.1 对 10.1.1.1 的 ICMP 报文。
- 接下来, client 回复该 ICMP 报文, 因此在 eth2 收到 ICMP 回复报文, 再向 eth0 转发该报文到 server1。
- 在前面接受 ARP 报文时 router 的 ARP 表学习了 client 和 server1 的 mac 地址, 因此第二次 ICMP 报文的请求和回复可以直接经由 router 进行转发, 不需要再发送 ARP 请求。
- 最后 eth2 中的 2 条 ARP 报文用于定期检查缓存。

Capturing from router-eth0

No.	Time	Source	Destination	Protocol	Length	Info
1	0.000000000	Private 00:00:00:00:00:00	Broadcast	ARP	42	Who has 192.168.100.2? Tell 192.168.100.1
2	0.047297229	40:00:00:00:00:01	Private 00:00:00:00:00:01	ARP	42	192.168.100.2 is at 40:00:00:00:00:01
3	0.047306877	192.168.100.1	10.1.1.1	ICMP	98	Echo (ping) request id=0x1303, seq=1/256, ttl=64 (reply in 4)
4	0.353128782	10.1.1.1	192.168.100.1	ICMP	98	Echo (ping) reply id=0x1303, seq=1/256, ttl=63 (request in 3)
5	1.000000832	192.168.100.1	10.1.1.1	ICMP	98	Echo (ping) request id=0x1303, seq=2/512, ttl=64 (reply in 6)
6	1.197042603	10.1.1.1	192.168.100.1	ICMP	98	Echo (ping) reply id=0x1303, seq=2/512, ttl=63 (request in 5)

Capturing from router-eth2

No.	Time	Source	Destination	Protocol	Length	Info
1	0.000000000	40:00:00:00:00:03	Broadcast	ARP	42	Who has 10.1.1.1? Tell 10.1.1.2
2	0.000029126	30:00:00:00:00:03	40:00:00:00:00:03	ARP	42	10.1.1.1 is at 30:00:00:00:00:03
3	0.105052129	192.168.100.1	10.1.1.1	ICMP	98	Echo (ping) request id=0x1303, seq=1/256, ttl=63 (reply in 4)
4	0.105078440	10.1.1.1	192.168.100.1	ICMP	98	Echo (ping) reply id=0x1303, seq=1/256, ttl=64 (request in 3)
5	0.948038095	192.168.100.1	10.1.1.1	ICMP	98	Echo (ping) request id=0x1303, seq=2/512, ttl=63 (reply in 6)
6	0.948065418	10.1.1.1	192.168.100.1	ICMP	98	Echo (ping) reply id=0x1303, seq=2/512, ttl=64 (request in 5)
7	5.351167200	30:00:00:00:00:01	40:00:00:00:00:03	ARP	42	Who has 10.1.1.2? Tell 10.1.1.1
8	5.432142320	40:00:00:00:00:03	30:00:00:00:00:01	ARP	42	10.1.1.2 is at 40:00:00:00:00:03